Validation and improvement of the SMPI simulation framework for MPI applications

by

Attila Döme Lehóczky

A thesis submitted in partial fulfillment for the
degree of Doctor of Philosophy

in the
Faculty Name
Department or School Name

2013. augusztus 11.

Declaration of Authorship

I, Attila Döme Lehóczky, declare that this thesis titled, 'Validation and improvement of the SMPI simulation framework for MPI applications' and the work presented in it are my own. I confirm that:

- This work was done wholly or mainly while in candidature for a research degree at this University.

- Where any part of this thesis has previously been submitted for a degree or any other qualification at this University or any other institution, this has been clearly stated.

- Where I have consulted the published work of others, this is always clearly attributed.

- Where I have quoted from the work of others, the source is always given. With the exception of such quotations, this thesis is entirely my own work.

- I have acknowledged all main sources of help.

- Where the thesis is based on work done by myself jointly with others, I have made clear exactly what was done by others and what I have contributed myself.

Signed:

_____

Date:

_____

„Write a funny quote here.”

If the quote is taken from someone, their name goes here

UNIVERSITY NAME (IN BLOCK CAPITALS)

Abstract

Faculty Name

Department or School Name

Doctor of Philosophy

by Attila Döme Lehóczky

Simulation has become increasingly popular in the scientific community lately because of its many advantages over using real systems in certain situations. Such an advantage is, for example, that we don't need to devote lots of money and other resources on acquiring the - possibly large - system we need for our experiments, or that we can easily switch between configurations, producing experiment results on more than one such system without a large overhead. SimGrid is a generic simulation framework, which can be used to simulate multiple kinds of distributed systems, such as Clouds, Grids or clusters. MPI is an inter-process communication protocol, one of the most prominently used in distributed computing currently. SMPI is a framework that connects SimGrid and MPI: it's part of the SimGrid project with the goal of simulating MPI applications in a simulated environment. SMPI is currently under heavy development. Creating generic network models is an ongoing effort and it is one of the development directions that gets the most focus by the development team. This process needs a lot of validation, which is done by running real-life MPI benchmarks (using possibly multiple different implementations), as well as running the same benchmarks with SMPI. The results can be used to check how accurate the simulation was. Currently, this testing process consists of multiple different steps. It's tedious, repetitive work, but as mentioned before, it's necessary. This thesis is about the implementation of a framework for automating such tests, introducing as much automation into the process as possible, providing the possibility to create test workflows that can be run with minimal user interaction. The goal is to speed up the testing process, making it possible to acquire proportionally more test results, thus helping the development team to make progress.

Acknowledgements

The acknowledgements and the people to thank go here, don't forget to include your project advisor...

Contents

List of Figures

List of Tables

Abbreviations

LAH   **L**ist **A**bbreviations **H**ere

Physical Constants

Speed of Light $\quad c \quad = \quad 2.997\ 924\ 58 \times 10^8$ ms^{-S} (exact)

Symbols

a	distance	m
P	power	W (Js^{-1})
ω	angular frequency	rads^{-1}

For/Dedicated to/To my. . .

1. fejezet

Introduction

Simulation has become increasingly popular in the scientific community lately because of its many advantages over using real systems in certain situations. Such an advantage is, for example, that we don't need to devote lots of money and other resources on acquiring the - possibly large - system we need for our experiments, or that we can easily switch between configurations, without having to worry about how to modify it or how to obtain a new one. SimGrid is a generic simulation framework, which can be used to simulate multiple kinds of distributed systems, such as Clouds, Grids or clusters. MPI is an inter-process communication protocol, one of the most prominently used in distributed computing currently. SMPI is a framework that connects SimGrid and MPI: it's part of the SimGrid project with the goal of running MPI applications in a simulated environment. SMPI is currently under heavy development. Creating generic network models is an ongoing effort and it is one of the development directions that gets the most focus by the development team. The development process needs continuous validation, which is done by running real-life MPI benchmarks (using possibly multiple different implementations), as well as running the same benchmarks with SMPI. The results can be used to check how accurate the simulation was. Currently, this testing process consists of multiple different, manually performed steps. It's tedious, repetitive work, but as mentioned before, it's necessary. This thesis is about the implementation of a framework for automating such tests, introducing as much automation into the process as possible, providing the possibility to create test workflows that can be run with minimal user interaction. The goal is to speed up the testing process, making it possible to acquire proportionally more test results, thus helping the development team to make progress.

Distributed computing is a very important concept in computer science. There are multiple types of large-scale distributed environments that can be used for either production purposes or for research. Parallelism is also used in a lower level, for example in graphics

processing, where we can have multiple graphics chips in one computer to do the task. When doing distributed computing, communication between the processes becomes a very important concern, since it can pose a relatively large overhead compared to sequential problem-solving - to make parallelism worthwhile, we have to make sure that the speedup provided by the distribution of tasks makes up for the communication overhead. To achieve this, task distribution needs to be carefully planned and the communication protocols are needed to be optimized. A widely utilized communication protocol that has been under development for many years is provided by the MPI[1] inter-process, language-independent communication API. MPI itself is just a specification, which has multiple existing implementations. The most widely used ones include OpenMPI[2] and MPICH[3].

Setting up a distributed environment is a complicated endeavour: it needs both human and monetary resources. When doing research on distributed computing, our needs possibly exceed the use of just one single platform: we would like to test how our experiments fare on multiple different environments. Such environments are not always at our disposal. This is why simulation has become a very important field of research. SimGrid[4] is a project providing a wide range of features in this regard: it is a scientific instrument that can be used to simulate large-scale distributed systems in order to study their behavior by evaluating and analyzing the results of parallel experiments run with the simulator. As mentioned before, inter-process communication is a very important concern for these experiments. SMPI is a framework that is part of the SimGrid project. This framework makes it possible to simulate the execution of parallel applications that use the MPI standard. This simulation can be done on a single node.

SMPI is a framework that has been validated in the past by experiment results. Such results have been documented and published, for example in [5]. Results are needed to be reproducible, by providing that the conducted experiments are repeatable. However, currently, the testing process to procure such results consists of multiple steps, many of them being manual configuration steps, such as the allocating of nodes, the creation of a file containing the allocated nodes or the distribution of the benchmark's runnable between the nodes. These tedious test processes lack a universal, user-friendly guide to help other researchers reproduce the acquired results in order to do further validation. Such a guide would also help in repurposing, extending our experiments by changing parameters, switching the underlying configuration, etc. Providing the possibility to easily repeat the experiments conducted for our paper is very important: our goal when writing a paper is not only to announce our results but also to convince our readers that our results are correct.[6] The best way to prove that we are right is if we provide a straightforward way for anyone to repeat our experiments and to see for themselves that our results stand. Also, if other scientists are unable tp repeat our experiments, that means that they are unable to profit from them: to incorporate our results in their research, which may or

may not be related to our field, which is the main point of scientific collaboration.

SMPI is a well-documented and working framework, but also an active project and as such, it is under constant development. Currently, a lot of resources are spent on developing generic network models for the simulator. Extensive testing is needed for the development process, since continuous validation is necessary to see whether or not we are heading in the right direction, whether or not the simulator correctly represents the real-world behavior. This testing process, as previously mentioned, is currently very time-consuming. This doesn't only limit the number of tests that can be run, but also limits the reproducibility of the results that are achieved with SMPI, the importance of which has been discussed in the previous paragraph. By constructing a framework that simplifies the testing process, more reliable and verifiable results could be produced, as well as it would make the SMPI project members' lives easier. It is important to include as much automation in the framework as possible, since other researchers that want to repeat the experiments might not be computer experts. Our processes can't be fully automated though. For example, certain parameters, such as the runnable, or the number of nodes to allocate have to be set - but we aim to keep such non-automated steps minimal.

This thesis discusses how such a framework could be built and provides an implementation, utilizing the XPflow[7] experimentation engine. XPflow is a fairly new project, constantly under development, but it's fairly stable and suitable for our project. It's designed to help automating experiments, utilizing a top-down approach taken from business workflows.

2. fejezet

Literature Review

In the literature review, we discuss the area of research covered by the thesis, citing relevant papers and articles that serve as a base of ideas for this document.

First, we talk about the Message-Passing Interface (MPI), the parallel programming API used for the purposes of this thesis. We also talk about the different implementations of this API, notably OpenMPI and MPICH. We discuss the methods of modelling and simulation in general. Then, we talk about SimGrid, which is a simulation-based framework, and SMPI, the implementation of MPI that runs on top of SimGrid. While discussing SMPI, we present the idea of a framework that would make it possible to automate running tests, trace collection and post-processing. In connection to this framework, we talk about workflows as a means to automate our experiments and whether scientific or business workflows are more suitable for us. We also talk about the importance of reproducible research, which is another argument for constructing such a framework.

2.1. MPI

Distributed computing is a very active and important subject of research in computer science, including fields such as cluster computing, grid computing, Cloud computing, or peer-to-peer computing. Communication between the different processes in a distributed application can be implemented in a number of ways. As communication is necessary in most cases, a standardized communication protocol can be a lot of help when developing a distributed program. The Message-Passing Interface (MPI) is a language-independent message-passing library interface specification. It is not a language, but a standard - there exist multiple MPI implementations. Since its take-off, it has become a de facto standard for inter-process communication. The standard provides vendors a clear set of

routines, that they can implement efficiently, or in a way that it suits the hardware they provide.[1]

2.1.1. OpenMPI

OpenMPI is an MPI implementation with the goal of being able to achieve good performance on a wide range of different aspects of high-performance computing. To efficiently support multiple types of parallel machines, high performance "drivers" for all established interconnects are developed. These include TCP/IP, shared memory, Myrinet, Quadrics, and Infiniband. Features for checking data integrity are provided in order to account for network transmission errors. With the utilization of message fragmentation and striping over multiple (potentially heterogeneous) network devices, OpenMPI provides an increased bandwidth to applications, as well as the ability to handle the failure of network devices during runtime.[2] On the Grid'5000 cluster, which is used to most of the research conducted for this thesis, OpenMPI is the default MPI implementation used by the default images.

2.1.2. MPICH

MPICH was originally developed during the MPI standards process starting in 1992 to provide feedback to the MPI Forum on implementation and usability issues. This original implementation was based on the Chameleon portability system to provide a light-weight implementation layer (hence the name MPICH from MPI over CHameleon). Around August 2001, development begun on a new implementation called MPICH2.[3] This implementation introduced improvements on collective communication operations by using multiple algorithms, choosing between them depending on certain variables - for example the message size.[8] Another important result during the development of MPICH2 was the design of the Nemesis communication subsystem and the porting of MPICH2 on that system. The efficient implementation of shared-memory communication helped Nemesis MPICH2 achieve low latency and high bandwidth.[9]
Starting with November 2012, the project is renamed to MPICH, with version number 3.0.[3]

2.2. Modelling and Simulation

In distributed computing, modelling means creating an abstraction of a real system by taking only the aspects of it that are relevant to the system's behavior into account. Once

constructed, such a model becomes a tool with which we can investigate the behavior of the system.[10]

2.2.1. Advantages of Modelling

Modelling and simulation techniques have been used extensively in parallel computing and is an ongoing research topic, with new challenges continuously arising. There are various reasons for its importance.

Conducting experiments on real-world systems can be infeasible because experimenting would disrupt the service that is provided by the system. For example, in the case of a mail server, experiments or monitoring could cause delay, or maybe even data loss. Service disruption can sometimes be even dangerous, in addition to being an inconvenience: in the case of a nuclear reactor, delay or loss of data can prove fatal. Timeliness can be as important in such systems as correctness. However, performance analysis and monitoring might be crucial to draw conclusions about maintenance, for example. Another problem with direct experimentation is that the information we are looking for may not be available, or may be complicated to get. For example, in most operating systems, it is difficult to obtain the exact timing of instruction-level events.[10] Also, when conducting experiments on a real-world system, results are often non-reproducible, due to resource dynamics.[4] Another argument on the side of modelling is that it provides the ability of experimenting on different configurations. Investing in a large-scale computer cluster, or the setup of a distributed grid environment is an expensive and tedious process. Investors want to make sure that they get what they want: they impose performance constraints on the system. This means that they want to know how the system will behave before buying it and setting it up. To predict the behavior, experiments are needed to be conducted. We need to do these experiments on different setups, before finding out which one is the best in the current situation. Changing the hardware or software configuration parameters on a real-world system is very inconvenient - in most cases, it's not doable, because of time and money constraints. Thus, the solution is to simulate the desired system, and run the experiments there. This way, changing the configuration is simple and costless.[10] Another great benefit of simulation is that in a classroom setting, students can learn the principles of high-performance and distributed computing without actual access to a parallel platform.[5]

2.2.2. Analytical and Simulation Models

The accuracy of a model can vary: we can make an analytical, or qualitative model, in which all definite values are abstracted away - in this case, we get a representation of

the system, which can be analysed mathematically to deduce its behavior. When using this method, no experiments can be conducted, we solely rely on theoretical analysis. In contrast, a simulation model is a stochastic model, which is an algorithmic abstraction of the real-world system that can be executed to reproduce the system's behavior. This model is also called a quantitative model, as we can get estimates of the modeled system's quantitative attributes, such as response time or throughput. In other words, we can use a simulation model to conduct performance analysis on a system, without actually having the actual system at our disposal.[10][11]

When wanting to get a prediction about how a specific system would perform, a theoretical model, in most cases, produces unreliable and unrealistic results - it's not feasible for such accurate predictions. The vast majority of research results are obtained via empirical evaluation of experiments.[4] For these reasons, we use the simulation model in this thesis. As we stated before, such a model can be executed, which is called simulation. During simulation, the model is supposed to behave like the real system would. It is hard to produce a 100% accurate simulation, but more and more reliable solutions are being developed. The simulation model contains more aspects of the real system compared to the theoretical model, in order to accurately represent the system, while still avoiding unnecessary detail.[10] Creating and executing a simulation model is complicated, computationally expensive and poses a number of challenges, thus, a good simulation framework (such as SMPI) can prove to be of much help when conducting experiments.

2.3. Off-line and partial on-line simulation

Full simulation - including CPU and network emulation - of a parallel application can be, in many cases, even more resource-intensive than running real-world experiments. This contradicts the fact that one of the most prominent goals of simulation is to observe the behavior of such large-scale platforms that aren't available. Thus, there is much interest in more efficient simulation approaches. [12] The most widely used of such approaches fall into two categories: off-line simulation, which is also called trace-based or post-mortem simulation and on-line simulation, which is simulation via direct execution.[5] As in the subject of this thesis, we are interested only in the simulation of MPI applications, we describe the two different simulation approaches concentrating specifically on that subject.

2.3.1. Off-line simulation

For conducting off-line simulation, logs or traces are needed to be collected of an execution of the MPI application to be simulated, taking place on a real-world platform. This is

necessary because the obtained traces are used as an input for the simulator, which then replays the execution traces as if the application was running on the target platform. This platform's characteristics may differ from the one's that we obtained the traces from, since we may want to use the simulator to predict the application's performance on a different system. Thus, there is a need to calculate how the target platform would execute the application, based on the traces we got on the other platform. The typical approach to this problem is to first compute the time intervals between the MPI communication operations. During these intervals, local computations were conducted, that's why we call these "CPU bursts". During simulation, we have to account for the differences between the performance of the platforms by modifying the time these CPU bursts take. This can be done by simply scaling the time intervals, or by using more sophisticated methods, by calculating exactly how the application's computational signature and the platform's hardware signature relate.[5] Communication operations, of course, also need to be simulated. This is done based on the events recorded on the trace, and on the network model of the simulated platform.[5]

As mentioned in [5], there are multiple downsides and challenges to the off-line approach. On such downside is that when wanting to simulate a relatively larger-scale application, the size of the obtained traces can be so large, that running the simulation on a single node might become a problem. Methods in order to overcome this obstacle include a compact representation of the traces in order to reduce its size. Another solution is to only consider a carefully selected subset of the obtained traces. A big disadvantage when using off-line simulation is that because we use the traces as an input to the simulator in order to replay the execution of the application, the simulation is dependant on the platform we collect the traces on. This means that, for example, there can be features in the obtained traces that might not be available on the target platform. In most cases, it is also necessary that the two platforms have the same number of nodes to run the experiment on. Although there has been a good amount of research done in the area, MPI itself and also the application might alter its behavior depending on problem and message size. Because of this, simulating the scaling of an application is a very hard, if not impossible task.[12]

2.3.1.1. Time-independent traces

Another link that ties the produced trace to the host platform occurs when we use timed traces, meaning that each traced event is associated to a time-stamp. Since the time delays between the events are specific to the platform specification, the simulator has to apply a correction factor to these delays when running the simulation on the target platform. Thus, the simulator has to know precisely the specification of both the host

and the target platform, in order to be able to calculate this correction factor. Another difficulty regarding that comes up regarding this problem is that actually calculating the correction factor is a tedious process. It can take a considerable amount of time, depending on how similar the host and the target platforms are.[13]

In [13], a solution to this problem is proposed: time-independent traces. Acquiring time-independent traces means that the traces won't contain any timestamps, breaking this link between the acquisition and the replay of the traces. In these type of traces, for each computation or communication operation, we log the volume of the operation (in number of floating point operations or bytes) instead of the time the execution took. This type of information, in most cases, does not vary depending on the platform the experiment is run on. The exceptions are the adaptive MPI applications that modify their execution path according to the execution platform.

[14] contains a guide describing how to acquire such traces on the Grid5000 platform. The guide was used to serve as a base for the process on how to acquire traces. Since the work in this thesis is mostly related to producing traces for validating on-line simulation, in which case time-stamps don't have an influence on the process (neither in a positive, nor a negative way), the extraction of time-independent information from the traces can be omitted in our case.

2.3.2. Partial on-line simulation

Partial on-line simulation is a different approach. Here, we execute the program with no or very little modification on a host platform, that tries to mimic the behavior of the target platform.[5] Computational tasks are executed on the hardware, but the timing and the delivery of the messages is calculated by the simulation environment. Thus, the simulator is responsible for mainaining the correct order of the events, both computational and communicational.[12]

A downside of the on-line approach is that since we actually execute the code, the resource needs for running the simulation is about as high or even higher (in case of needing an extra node to run the simulation component, for example) than it is for the actual experiment. Techniques have been implemented in order to help alleviate this problem. The basic idea is that the actual results of the experiments (for example, the result of multiplying two matrices) might not be important in our case: we are only interested in the *time* it takes to get those results on the target platform. This is why methods can be employed which trade off accuracy for performance. This idea might not be feasible for experiments where data-dependent application behavior is vital, but a large portion of benchmarks can be indeed simulated this way, providing a reasonably accurate execution profile.

Although slower, on-line simulation is more general than the off-line approach, as it does not, in any way depend on some other platform - whereas in the case of off-line simulation, as we mentioned before, the trace is acquired on a different platform, with maybe specific application configurations, thus inevitably bringing dependencies.

2.4. SimGrid

For reasons mentioned before, simulation techniques have historically been widely utilised in several areas of computer science, e.g. microprocessor design, network protocol design. Due to this, a lot of effort went into developing the technology and as a result, widely used and reliable simulation frameworks have been developed in these areas. However, there hasn't been a well-developed standard simulation tool for what we talk about in this thesis: execution of distributed applications on distributed computing platforms. Rather, there has only been a number of in-house developed, highly specialized tools to satisfy the need of the community. SimGrid is a more generic simulation framework that is being developed to be one of the acknowledged and widespread tools for simulation in large-scale distributed computing.[4]

SimGrid's key features include:[4]

- A scalable and extensible simulation engine that implements several validated simulation models, and that makes it possible to simulate arbitrary network topologies, dynamic computational and network resource availabilities, as well as resource failures;

- High-level user interfaces for researchers (who are not necessarily computer science experts, but rather experts on their own field of research) to quickly assemble simulation prototypes in either C or Java;

- APIs for distributed computing developers to create distributed applications that can run seamlessly in either "simulation mode" or "real-world mode", in order to be able to test it on the simulated environment before actually deploying it.

SimGrid is a very active project, both in terms of research and in terms of development. It is a favored tool by researchers, which is proven by the increasing number of papers written where the research was conducted using SimGrid as a scientific instrument. In terms of development, the developer team envisions a number of directions for future work: addition of a model for disk resources; extention of scalability to improve usability in the P2P domain; ability to dispatch simulated nodes over several physical machines.[4] Another important field of research for the SimGrid team is the implementation of the API that has already been mentioned: the Message-Passing Interface (MPI).

2.5. SMPI

As stated before, MPI is one of the most widely used APIs for communication between nodes in distributed computing. SMPI is a framework for simulating on a single node the execution of parallel applications implemented using the MPI standard. It is part of the SimGrid project and as such, it is built on the SimGrid simulation kernel, benefiting from its fast, scalable and validated network models. SMPI also extends the existing model with other techniques, such as a validated piece-wise linear model for data transfer times between cluster nodes. SMPI simulations also account for network contention - timing and delivery of the messages are determined using the network model of SimGrid.[5] A current limitation in SMPI is that it is unable to simulate high-performance networking hardware such as Infiniband. Thus, when wanting to compare simulation to real-life results, we have to make sure those results were gatheres using Gigabit ethernet.
Three of the main challenges for simulating an MPI application are:

- Accuracy: The prediction of the real-world execution time (the "simulated time") needs to be as accurate as possible, so that reasonable conclusions can be drawn from the experiments.

- Scalability: We want to be able to simulate large-scale applications within a reasonable timescale.

- Speed: It would be advantageous if the simulation time (the actual time of running the simulation) would be as low as possible, compared to the simulated time (the predicted execution time of the real-world application).

As for simulation methods, SMPI can be used for both off-line and on-line simulation, although the emphasis is more on the on-line approach, since it's actually a partial implementation of the MPI standard in itself, thus making it feasible for executing MPI experiments. More specifically, in SMPI, the goal is to be able to make such simulations on a single node. The most prominent challenges when doing this are the large CPU and memory requirements. SMPI provides some special techniques that help overcoming these challenges. The basic idea about trading off accuracy for performance has already been described in the previous section about on-line simulation. SMPI implements multiple such techniques, allowing to run experiments with such high resource requirements that would otherwise be impossible to fulfill. Such a method in order to reduce CPU usage is to run the benchmark only on a subset of all the nodes, while in place of running the code on the others as well, we just insert the computation time that we got previously. Apart from CPU usage, we need to also account for the need for memory. A technique for that is "RAM folding": here, multiple simulated processes, that in SMPI

are, in fact, simple threads, use the same reserved memory location, thus overwriting each other's data structures. Also, another implemented solution is to remove large data array references from the code, with the help of the compiler which can result in the complete removal of potentially large, now unreferenced arrays. Again, this obviously corrupts the results that the experiment program gives, but in the same time helps to simulate applications that would use such an amount of memory that just wouldn't be physically possible to provide in our testing environment, while still providing a reliable estimate of the performance.[12] These features are disabled by default, they have to be explicitly enabled by the user.

Extensive testing was conducted in [5] to verify the previously mentioned qualities of the framework. In these tests, the OpenMPI and MPICH implementations were used to serve as verification benchmarks: the same experiments were run using both MPI implementations, as well as simulated with SMPI. The results show that SMPI predicted the execution time of OpenMPI and MPICH applications for point-to-point, one-to-many and many-to-many applications with an average error value of under 10% in each cases. Using the aforementioned techniques to reduce the memory footprint, SMPI tests were successfully conducted on a scale of up to 448 processors. The results showed that the predicted execution times were underestimates with an average error value of 18.5%, which is higher than in previous experiments without these techniques. We have to note here, though, that certain tests weren't successful without the RAM-folding techniques, due to an out-of-memory error. This shows, that although it poses difficulties, reducing the memory usage is vital in SMPI.

As SMPI is an actively developed project alongside SimGrid, there are a number of research directions. One major development to the project would be a testing framework that would aim to lessen the burdens of testing as much as possible. The goal is to provide a unified method to set up experiments across different environments and to do it with as little necessary adjustments on user part as possible.

2.6.   Experimentation tools

Experiment orchestration and process automation is not a new idea. There exist multiple tools for doing this. Among others, such tools are Expo[15], Plush, OMF[16] or the Grid5000-specific g5k-campaign. The problem is that most of these tools use a bottom-up design, meaning that in order to understand and use the experiments, the user has to understand the lower-level building blocks first. A top-down approach would make it possible to start the design of the experiment with a high-level description, then work our way down to the lower level details while implementing. There already exists an

approach like this, in Business Process Management (BPM).[7] Before talking about BPM though, let's take a glance at workflows in general.

2.7. Workflows

When talking about a framework to automate MPI experiments, we are essentially talking about creating workflows: we connect multiple steps, making it possible to execute them in a chain, with no or minimal user interaction during the process. In the application level, workflows are abstract in the sense that the workflow only describes the goal of the experiment, its components and its dependencies. Lower level implementation details are hidden from the user of the workflow. This provides the possibility of changing the implementation without having to change the high-level workflow description - as long as the new implementation still has the properties that are written in the description.[17] In the subject of workflows, there are two main approaches that we will discuss: scientific workflows and the previously mentioned business workflows.

2.7.1. Scientific workflows

Scientific analyses often have to be conducted in several different hardware and software environments. Exporting and importing data from and to different environments can be a tedious task, slowing down the work process, forcing researchers to divide their efforts between computation management and their actual research. This is the main reason scientific workflows are widely used in various different scientific domains: they are a formalization of the ad-hoc process that a scientist has to go through in order to get from raw data to publishable results. Since the raw data to be analyzed can be large, heterogeneous and complex, the process can be computationally intensive and produce derived data formats, which is one of the main differences between scientific and business workflows.[18]
There are several tools that help in experiment design, mapping of computing resources to the workflow and handling exceptional situations. Some of the more well-known tools include Kepler[18], Pegasus[17] and Taverna[19].
Scientific workflows are well suited for managing computation on *a priori* available data or data queried from public databases. However, it's not well suited to cases when data acquisition is actually part of the workflow process, or in other words, when the source of the data is the computer system itself. They are also data flow-driven, which is not true in the case of our processes that we want to automate. This makes scientific workflows not optimal for research conducted in computer science.

2.7.2. Business workflows

Business workflow management systems are usually based on agreed-upon standards in order to facilitate communication between different software systems and companies. The workflow logic is control flow-driven and includes constructs to specify paths and conditions.[20] The top-down approach it uses is what can make it viable in our case: in Business Process Management, the first step is to understand the organization. Then we can model its processes as workflows and execute and monitor them. While monitoring, we can spot defects and work out ways to improve the organizational activities, as well as we can redesign the processes to make them cheaper and faster.[7]

This approach can be utilized in the domain of experiment orchestration, making business workflows a viable choice when approaching our problem. An experimentation engine with the goal of implementing this idea is XPflow[7]. This engine was used to implement the test automation framework in this thesis and will be discussed in a bit more detail later on in this document.

2.8. Reproducible research

New scientific ideas, developments and results are only useful when they are documented and published. It is vital that results are announced, so others can be aware of the latest developments on their field of research. This helps in creating a linked data cloud, used by scientists to incorporate various output of other research into their own, using previous results as "stepping stones" to achieve something new.[21] But simply publishing results is not enough in order for others to make use of them. Besides announcing the achievements, the other goal of scientific publications is to convince the readers that the results it presents are correct. Besides theoretical reasoning, papers in experimental science should provide a documented methodology describing how the author has gotten to those results.[6] The methodology has to be detailed and precise enough so other researchers can repeat the same steps, thus reproducing the same results. This is vital in order to provide the possibility to verify those results and to fully understand them. We have to note that in a reproduced experiment, it's not the *raw results* that need to be identical: they merely need to fit within a statistical margin, compared to the original results, so that the *conclusions* derived from them can be the same.

Reproducing the results also makes for a starting point for further development, as the described methods used for reproduction can be extended to achieve something more or something different in the same area of research, or repurposed to gain useful results in a completely different area. Reproducibility is a relevant concern in the case of SMPI and one of the main goals of the testing and validation framework is to make developments

in the area.

In this chapter, we reviewed the background material to the general concepts related to this thesis. We talked about the MPI inter-process communication API and its two different implementations: OpenMPI and MPICH. Then we talked about modeling and simulation in general and how simulation is advantageous in certain situations. Relating to simulation, we talked about SimGrid, a multi-purpose simulation framework. Then we talked about the project that interconnects these concepts: SMPI, which is a tool that is able to simulate MPI programs on SimGrid. We talked about its features and how it's been previously validated with extensive testing. Since testing is currently a tedious multi-step process and continuous validation is necessary, the project is in need of a test automation framework to help with the development process. Also, such a framework would make documented experiments repeatable, as well as the results reproducible, which is an important concern as well. The subject of this framework is connected to the subject of workflows, which we discussed in some detail in the section after. We discussed workflows in general, as well as scientific and business workflows. We concluded that business workflows are more suitable for our current problem because of its top-down, control flow-driven approach. We briefly mentioned the XPflow experimentation engine, which is based on this approach and is utilized in implementing the framework. In the end, we also discussed the importance of reproducible research, which is another important aspect in our motivation for wanting to create the automation framework.

3. fejezet

Problem Description

In this chapter, we approach the problem regarding the creation of our framework from different aspects. First, we describe the most important attributes we'd like the framework to have. After that, we talk about how trace collection is done, elaborating on the current manual process of obtaining real-life (RL) traces, illustrating why it's repetitive error-prone, thus emphasizing the need of automation. After that, we shed light on some obstacles regarding tracing, such as tracing overhead or clock synchronization problems. While discussing these, we also mention some possible ways of overcoming these obstacles.

3.1. Testing framework

As we previously discussed, there are two main reasons why a test automation framework is needed by the SMPI project. It's a fact that the current process is undocumented, long, repetitive and error-prone. Thus,

- the heavy development that is currently ongoing in the SMPI project is impeded, since development brings with it a constant need for validation, which is done by running test.

- The effort required to reproduce research results related to SMPI are needed to be minimized.

When specifying the framework, we need to lay down a certain set of desired properties that we want it to have. The following is a brief description of these properties, relying on some of the ideas in [22]. Some of these attributes will be elaborated on in greater detail later on in this document.

Scalability.   We want to create experiments with this framework where if we increase the amount of resources (for example, we run the experiment on more nodes), then there is no unnecessarily drastic increase in the preparation or the execution time, the results remain correct and the execution is free of failures. The unavoidable errors need to be handled accordingly, providing the user with the necessary details of what went wrong.

Descriptiveness.   We want to describe our experiments in a top-down approach. This means that first, we describe the meaning of the experiment: what its goal is. The lower-level details (how the experiment is conducted, what tools we use, etc.) are also provided, but they are not necessary to understand the purpose of the experiment. Providing the possibility to ignore the details helps to quickly understand the different experiments, even for newcomers.

Modularity.   We want our experiments to be built from independent, interchangeable "building blocks". This way, existing experiments can be used as a starting point: certain parts can be modified, while others can be left in place, creating a workflow that achieves something different. Also, modularity helps maintaining the framework: the different parts can be modified, improved, or even completely rewritten, not affecting the other modules in the tool.

Reusability.   This attribute intertwines with *modularity* in the sense that we want to be able to reuse certain parts of the experiments when creating new ones. This not just eases the work of others in the research group, but also provides any researcher the possibility to reuse already written code, so they can reproduce our results or modify the code to alter the behavior of the workflow.

Error handling.   As mentioned in *scalability*, there are unavoidable errors that happen from time to time: network issues, unreachable nodes, deployment failure, etc. Such errors need to be handled gracefully: the program needs to provide information to the user as of why the experiment is failing, or why the runtime is prolonged. Whenever possible, an attempt to somehow manage the situation has to be made before exiting. Such attempt can be multiple retrying, allocation of a different set of nodes, etc.

Minimal human interaction.   The main goal of this framework is experiment automation. However, not all human interaction can be avoided: certain parameters are needed to be provided (e.g. the number of nodes, the desired site to use, etc.), or manual assessment of the results may be needed to be conducted. That said, the goal of the

framework is to limit the amount of human interaction required to what is absolutely required to run the actual experiment.

Metadata collection. The usefulness of the framework can be greatly elevated if it provides the possibility of some (even minimal) metadata to be collected about the experiment that was run. Such metadata can include the description of the goal of the experiment (provided by the researcher), the date the experiment was run, the nodes and the operating system images that were used. The metadata can be collected and stored in a specific location (for example in a database). After a large amount of experiments were run, such a collection of data can be queried to search for specific results. For example, the results of a specific experiments can be collected, compared and maybe visualized, or queries about a certain set of nodes can be made in order to find faulty ones.

The above list doesn't necessarily include all the technical properties of the framework, but it's a list about the most important design standpoints taken into consideration. After the bried description of these points, we talk about the actual process that we are aiming to automate.

3.2. Obtaining traces

SMPI is an actively developed project and as such, a lot of tests are run and a lot of measurements are taken. Previous papers ([5] [12]) have shown, amongst other results, how accurate the performance predictions SMPI makes are and how the time of the simulation can be lowered, while getting very little differences in simulated time (which means the predicted performance). These results are obtained through extensive testing and as development continues, more and more test data is needed for verification purposes. In order to get results, both real-life (RL) tests and on-line simulation tests using SimGrid (SG) are needed to be obtained and then the results need to be visualized, analyzed and compared. Comparison is very important, since this is how we can validate SMPI. All this needs to be done in as many different kinds of environments as possible. Currently, this is a tedious task that needs a lot of configuration by hand, as there is no unified methodology or automation for setting up experiments and obtaining results on different kinds of distributed environments and for different kinds of MPI implementations - one has to find his/her own way to make it work. Documentation only exists for specific systems, for example in [14], there is a guide about how to produce time-independent RL traces on the Grid'5000 testbed. Obviously, platform-specific guides don't always help with problems arising in an other environment.

Below comes a more detailed description of the process of acquiring traces. We're only talking about collecting traces of real MPI experiment runs, so-called real-life (or RL) traces, since this is our current main concern with the thesis. Collecting traces of simulation runs is a different process, which could also be integrated into the developed framework in the future.

3.2.1. Real-life (RL) traces

As talked about in detail in [14], conducting RL tests involves multiple steps. RL test data collection is done by collecting traces of MPI benchmarks. Currently, the favored tool in trace collection in the project is Tuning and Analisys Utilities (TAU)[23], which is a well-established profiling tool, that also provides tracing features. Thus, on the system where we are running the benchmarks, TAU has to be deployed and configured, alongside other software that TAU depends on. One is PAPI[24][25], an interface which provides us with the possibility to get access to low-level hardware counters (to trace the number of instructions at processor level). We also need the Program Database Toolkit (PDT)[26], which provides the ability of automatic performance instrumentation. In order for TAU to collect the traces we need, these toolkits have to be deployed and correctly linked with TAU. TAU has its own compiler scripts for MPI programs for both Fortran, C and C++. After compiling a benchmark using one of these scripts, they will generate TAU trace files upon execution. One trace file (.trc) and one event file (.edf) is generated for each MPI process. The traces are generated on the node of execution, that's why we need the trace_gather program[14], which is an MPI program that gathers all the trace files on one node.

It is possible to visualize the traces, in order to compare the RL and SG results more easily. Paje is a visualization tool that can be used for this purpose. It has its own trace format that it can comprehend, thus the TAU traces have to be converted to that. Also, we need to merge the traces into one file that we can give to Paje after the conversion. There is a TAU script that is able to do this, creating one trace and one event file. We now only have the task of converting the TAU trace to a Paje trace file. Another MPI tracing library, Akypuera[35] provides this possibility, having its own tau-to-paje conversion script. Once done, we finally have a trace file that Paje can read and display to us.

As we can see, this process is made up of multiple steps. It's repetitive and prone to errors when done manually, which causes frustration and unnecessarily delays when testing. In the future, it is very likely that more and more tests will have to be run in order to verify old and newly implemented features, as well as various experiments will be conducted using SMPI to test it against various MPI platforms. The main reason of

importance of developing an automated way for testing and validation is that it could make the previously discussed tedious trace-gathering process a lot smoother and faster, with less user interaction, thus less delay and frustration. If the process could be incorporated into a single workflow, proportionally more tests could be run, providing more reliable results with less effort than before. Apart from making extensive testing and experimentation more straightforward, a functioning test and validation framework would provide a documented method, which could be used by other researchers to reproduce the achieved results, the importance of which has already been discussed above.

3.3. Tracing-related problems

Apart from wanting to simplify the previously described, fairly convoluted process of getting results, there are other problems related to tracing that need to be addressed. Since we need to compare RL and SG traces, it is very important that the traces we get are accurate, otherwise we could be lead to false conclusions. Below are the most prominent traps that could potentially compromise our results. Note that these are problems that come up in a real environment, thus only related to real-life traces.

3.3.1. Multiple cores on one node

Nowadays, it is common for a computer to have multiple processors (cores). Because of this, applications have been optimized to increase performance by utilizing more than one core if possible. MPI implementations have also been optimized to migrate threads between cores in certain cases. The problem with this is that SMPI is not configured to account for the speedup that the utilization of multi-core processors can cause. So if we compare the running time of an MPI application that was run on multi-core nodes and compare the running time to our simulation's running time that was done by using SMPI, even if we simulated the used platform correctly, we will likely see that SMPI overestimated the running time, since it didn't take into account the performance increase caused by using multiple cores.

A solution to this problem is to explicitly disable all cores but one on every node before running the real-life experiment, as seen on figure 3.1.[14] This can be done since we only want to run one MPI process on each node, thus the other cores won't be utilized by other processes in our experiment. The downsides to doing this are that we need to know the platform-specific instructions on the nodes which have to be given to disable cores, as well as we most likely need root privileges. But if we can use it, this method is a simplistic and sure way to solve this problem.

Another possible way to make sure that MPI doesn't make use of having multiple cores

at its disposal is to use process binding. For most MPI implementations, processes can be bound to nodes (meaning that processes can migrate freely among the cores of the node), sockets (processes can only migrate between cores in the same processor socket) or cores (that is, a process is bound to one core, migration is disabled). In our case, we need the latter option: binding processes to cores. The difference between disabling the cores and this method is that here, none of the cores are disabled - with this option, we simply tell MPI to not use migration (figure 3.2). In OpenMPI 1.4 and above, this can be done by using the processor affinity instruction parameter *--bind-to-core* when running the application. In MPICH, the option *-binding cores* can be used to enable this feature.

Figure 3.1. One core is hosting the MPI process, others are disabled

Figure 3.2. None of the cores are disabled, but the process is bound to one core, so it doesn't migrate

3.3.2. Impact of instrumentation

In the context of the collection of traces, "instrumenting" an application means specifying what kind of data we need and which part of the program we need it from. This can be done either directly (by inserting function calls or macros that record the data we want), or by using an overlay library for this purpose. For example, in the case of TAU,

there is a feature called selective instrumentation, which provides us with the possibility to specify which functions we want to be traced.

When instrumenting the application, it is important to know at what degree we want to do it. We want to collect the data we need, but the greater the degree of performance instrumentation in the program, the higher the likelihood that the performance measurements will alter the program's performance, an outcome called *performance perturbation*.[23] In the case of most performance tools, TAU included, this is a concern that the developers try to address by reducing the overhead of the performance measurements as much as possible. It is worth noting though that although the overhead the measurements cause might be reduced, they can't be completely eliminated, since the tracing operations have to be handled by the processor as well. Because of this, the user has to be very careful when instrumenting an application. There are two main concerns that come up in our case.

Time overhead. Instrumentation can have a sizeable negative impact on the performance of the application. If the instrumentation causes a noticeable increase in running time, the simulator's prediction might be off, since it doesn't take into account the instrumentation overhead (as instrumentation is not part of the experiment).

Impact on hardware counter values. As mentioned before, when collecting time-independent traces, we use hardware counters to measure the volume of each operation. The hardware counter doesn't distinguish between events related to the experiment and tracing operations, thus all of them are taken into account. The result is that the collected trace represents the traced experiment, while we want information from just the experiment, without the traces. If the difference is too big, this can make our simulations look inaccurate. The impact it can have when using off-line simulation, where the simulator replays the traces corrupted with this overhead is shown in [27].

In [27], the authors propose an instrumentation as a correction to their previous work in [13]. The problem was that there was a sizeable overhead caused mainly by TAU building the whole call path of the instrumented application. While it can be very useful when trying to find bottlenecks in the application, this information is not needed for simulation purposes. In the proposed method, we tell TAU to exclude all of the source files from the instrumentation. This way, instrumentation becomes minimal, while still covering our specific needs: the hardware counter will be triggered at each MPI call to measure the number of executed instructions in the operations. All of the information related to MPI calls, i.e. the id of the process that made the call, the name of the function and its

parameters are traced. All this while both the time overhead and the hardware counter value discrepancies are considerably reduced, as shown in [27].

3.3.3. Clock synchronization

Another obstacle that comes up when wanting to analyze trace data generated in a distributed system is related to clock synchronization. When running a parallel experiment on a distributed system that uses multiple nodes, all the used nodes produce separate trace streams independently of each other. The problem is that between the local processor clocks of different nodes, there almost always exists some amount of discrepancy, mostly related to the temperature of the processor. No matter how little this discrepancy is, it can accumulate over time, as well as it can change the logical event order, which requires a message to be received only after being sent from the other node. This is also referred to as the *clock condition*.[28][29] Such inaccuracy in the traces can lead to false conclusions when doing the performance analysis. Even though in our work we want to use time-independent traces, violations of the clock condition is still a problem we have to be able to handle.

The problem could be easily avoided if every process would use a global clock instead of the local clock of the node it's running on. The problem with this approach is that accessing the global clock can be much more expensive than accessing the local clock, thus causing performance issues, as well as it's not available on every platform. In [28], the authors use an IBM switch adapter's globally synchronized clock's register to periodically collect global clock records for each node, thus being able to correct the local clock discrepancies after the experiment finished.

The main problem with this is, as already mentioned, that although some systems do offer a relatively accurate global clock, many other systems are only equipped with processor-local clocks, in which case the problem has to be approached from another direction. There exist clock synchronization protocols, such as NTP [30], which provides a widely used solution to align the clocks to a certain degree by adjusting local clocks at regular intervals to a globally accessible time server. Unfortunately, due to varying network latencies, this method still leaves an error rate of about 1 ms when synchronizing, which is not good enough for our purposes. TAU handles this problem with another post-processing approach, using an extended and parallelized version of the *controlled logical clock* (CLC) algorithm, which is described in [29]. CLC retroactively corrects clock condition violations in event traces of message-passing applications by shifting message events in time. When making such a correction, CLC makes certain precautions, since after the modification of individual timestamps, the length of intervals between events in the immediate vicinity of the affected event might change, as well as new clock violations

might be introduced.

In this chapter, we talked about the reasoning behind why we need a trace automation framework. We described certain attributes that we want our tool to have, as well as the current trace collection method that we'd like to make easier and more straightforward. We also depicted some of the obstacles that we need to pay attention to when collecting traces, as well as methods to overcome them. In the next chapter, we'll discuss more specific details about the implementation of the framework.

4. fejezet

Implementation

In this chapter, we discuss specific details about the implementation of the framework. First, we talk about testbeds, specifically Grid'5000, which is the testbed that was used during development. We discuss the main features of Grid'5000, as well as some important usage details. The first version of the test framework is Grid'5000-specific, meaning that it utilizes features that might be different or might not be present on other platforms. In the long term, the framework can be developed to eliminate such dependencies. After that, an example process is described, that represents a typical use case of our framework: an MPI experiment on multiple nodes. We also discuss XPflow, the experimentation engine used for the implementation. We talk about XPflow's approach to creating workflows, taken from Business Process Management. We describes how this is applied to our example experiment by depicting the previously mentioned example process in an XPflow model, consisting of *processes* and *activities*.

4.1. Environment

4.1.1. The importance of testbeds

Most Grids that are deployed at a large scale are production platforms inappropriate for research: such Grids mostly have an environment that's been set up specifically for the owner's purposes. Such a system would most probably need some amount of reconfiguration in order to make it feasible for what we would like to do, which might influence the behavior of the system. Another concern is that running experiments on the platform might cause delays, or even disruption in the service the Grid is originally used for. Obviously, this is most probably not acceptable for the owner of the Grid. This is why it is important to make a distinction between production Grids and Grids

that are made for testing purposes, or "testbeds". Because testbeds are specifically made for researchers to run experiments on, using up resources is not that much of a concern as it would be on a production platform. Since other people might be using the same platform, there are of course still some limits as to how much resources one user can utilize, but these limits are not so strict and are much more prone to negotiation. As previously mentioned, another important factor is that the nodes we are working on should be reconfigurable: we need to be able to make customizations to set up our testing environment. We need to be able to install and uninstall programs. Root access should not be necessary when doing tests, but it can make things easier. Deep control and monitoring mechanisms are also needed (not just in one node, but across multiple nodes) in order to be able to track our experiments.

Simulation with SMPI can be done on any system, since it only needs one node but in order to run RL experiments, such a testbed is needed. Most of the work related to this thesis has been done on the Grid'5000 platform.[31]

4.1.2. Architecture of Grid'5000

Below, we discuss some of the architecture aspects of Grid'5000 to show how it fulfills our previously mentioned needs, as well as how it addresses some other concerns as well. Description details are taken from [31].

4.1.2.1. Networking

Grid'5000 is a platform currently consisting of 5000 CPUs, distributed across 9 different sites in France (figure 4.1), connected by high speed network. These sites host their own clusters and they are connected through the Internet. It is very important with regard to the experiments that inter-site communication and inter-node communication are unrestricted and don't weigh any overhead on the experiments. Thus, all communication can be done without any constraints between sites. But as for security, we have to take into account the following: if a node is fully reconfigurable by the researcher, that means we can't make any assumptions about the configuration of the security mechanisms on an allocated node, thus, we have to assume that they are unprotected. This is the reasoning behind the decision that sites themselves (thus, of course, the nodes they are hosting) are not directly connected to the Internet, making Grid'5000 an isolated domain. This way, the sites are protected against DoS attacks.

It is possible to open restricted routes through the Internet to external clusters, which provides the possibility of doing multiplatform experiments.

Figure 4.1. The sites of Grid'5000

4.1.2.2. User view and data management

As previously mentioned, communication is done with minimal restrictions between Grid'5000 machines, meaning that authentication procedures in such cases is also minimal: a user has a single account across the whole platform. An LDAP directory is installed to provide this in a reliable way. Every site runs an LDAP server. These servers have the same root and there is a branch for every site. On a given site, the local administrator has read-write access, as well as is able to manage user accounts.

A user has access to all of the Grid'5000 sites and services (monitoring tools, wiki, deployment, etc.). The user also has an independent home directory at every site as well. Synchronization can be done with any of the standard tools, such as *rsync*, *scp*, or *sftp*. Data transfers to the outside world are restricted - it can only be done via secure tools such as *scp* in order to prevent identity spoofing. Authentication is done via a user-generated public key, in order to prevent brute-force attacks.

4.1.2.3. Experiment scheduling

At cluster level, the OAR[32] resource management system is used to handle the scheduling of experiments and resource allocation. Large-scale operations such as parallel task launching or monitoring is handled by a specialized parallel launching tool, Taktuk[33]. Taktuk is a handy tool which can be used from the console to, for example, perform a certain set of tasks on multiple nodes.

A simple grid broker is handling resource management at grid level, allowing co-allocation of nodes on multiple clusters. Co-allocation is a very simple process for the user who, after submitting an experiment needing several sets of nodes across different clusters, receives an identifier from the broker which can be used to retrieve all necessary information about the allocated nodes.

Node reconfiguration, talked about in more detail below, co-operates with the resource management system at certain points. Such a point is that when a user submits an experiment that requires node reconfiguration, the job submission is registered in a queue. Also, in the prologue script that runs before the actual experiment, deployment rights are given to the user which gives him/her the capability to deploy system images on the allocated nodes. An epilogue script runs after the experiment, revoking these rights. After the experiment is finished, all of the allocated nodes are rebooted, deploying a default environment, to provide a constant, unified system to run experiments on that don't need node reconfiguration.

4.1.2.4.   Node reconfiguration

Node reconfiguration on Grid'5000 is handled in a very user-friendly way, using a tool called Kadeploy3[34]. For every user, a set of default environments is available at start. After starting an interactive job (that is, a job that requires node reconfiguration), the user can deploy any of these environments by providing the chosen image's name to Kadeploy3. Deployment usually takes only a few minutes to complete - deployment time increases if we do the deployment on more nodes. After this, the nodes are rebooted and the user can log onto any of the nodes where an image was successfully deployed. When logged in, the user can freely customize the environment: he/she can install or remove software, modify configuration files, etc. - root password is given to Grid'5000 users as well to provide more possibilities. After reconfiguring the environment, the user has the the possibility of saving the now customized image. This image includes all software layers from OS to application levels, just as it was for the default environments. The home directory is independent of the image. After successfully saving an image, the user can deploy it on the allocated nodes the same way he/she did for the default images. This way, an environment tailored for the specific needs of the user only needs to be created once, then it can be freely reused on any other node on the site at a later time. When trying to port an image created on one site to another site, there can be compatibility issues due to inter-site differences. Modifications to the customized image can be done with ease.

4.2.   An example process

To reiterate the more detailed description of trace acquisition, an example process that we would like to automate consists of the following steps, assuming that we already have a customized image set up for our tests, as well as we already have an executable benchmark application we'd like to run:

- allocate the specified number of nodes on a specified site

- deploy our custom image on the allocated nodes

- create a nodefile containing the names of the allocated nodes

- broadcast the runnable across the nodes

- disable all cores but one on every allocated node (see 3.4.1 for explanation)

- run the mpi experiment from a chosen "head" node (can be any of the allocated nodes)

- gather the traces from the allocated nodes to the head node

- merge the traces

- convert merged trace file to Paje format

As a side note: in this example, we assume that we are working on the Grid'5000 testbed. As previously mentioned, most of the work regarding this thesis was conducted there.
As for the example process: parameters that the user can give to it are the number of nodes, the chosen Grid'5000 site, the image to deploy, the runnable and parameters to mpirun (for example to disable Infiniband) and the trace_gather script. The nodefile created in the 3rd step is necessary for the operations that are needed to do something on all the allocated nodes.
Now let's take a look at the experimentation engine we'd like to use for implementing the framework to automate processes like this.

4.3.   XPflow

XPflow is an experimentation engine, with one of its main goals being to provide the possibility to easily automate experiments by creating workflows. XPflow is a fairly new project. So new in fact, that at the time of writing this thesis, its source code hasn't been published yet, although it's going to be released in the very near future. In 2.8.2,

we already discussed business workflows, foreshadowing the fact that XPflow, the tool used to implement our framework, is based on that approach, which includes Business Process Modeling and Management. There are 2 main concepts in XPflow[7]:

- Process: It is the high-level description of the experiment, written in a DSL, which is embedded in Ruby. Processes are responsible for orchestrating activities and other processes, creating a workflow.

- Activities: The low-level building blocks of the experiments. They are written in Ruby and used for implementing the lower-level details of the experiment, to do the "real work" in it (for example: manage files, start the MPI job).

Lets take a look at how these concepts are used through examples, while uncovering a few other possibilities of XPflow.

4.3.1. A general example

First, we examine an example of transforming general, everyday activities transformed into XPflow (4.2). On the diagram, circles represent the starting and end point of the
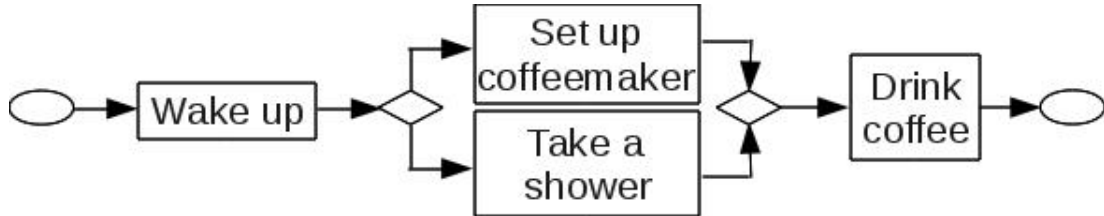

Figure 4.2. Morning routine in XPflow

process, the rectangles represent activities and the whole diagram represents a process. In XPflow, it is possible to execute activities sequentially or in parallel. An example for sequential execution on 4.2 is drinking coffee: the activities preceding it must all be finished before that activity can be started. On the contrary, setting up the coffeemaker and taking a shower can be executed in parallel, since it's not necessary to stand by the coffeemaker while it finishes - we can take a shower while it's running.

4.3.2. Real-life trace collection example

After the generic introduction, let's take a look at how we could represent the experiment we described in section 4.2 in XPflow. The representation can be seen on 4.3. On this representation, there is no parallel execution like in the previous example: one

Figure 4.3. A typical MPI experiment in XPflow

activity has to be executed sequentially after the other, the activities depend on each other. However, as we can see, some of the activities are distinguished by the text "forall(nodes)" above them. This means that we iterate through all the allocated nodes, executing the activity on all of them. This is done in parallel.

As a side note: in the current implementation, this kind of iteration is realized not specifically by XPflow, but by the parallel launching tool mentioned in 4.1.2.3, Taktuk[33]. This is simply because it was simpler to implement it that way, because this is a known and stable solution, while XPflow is under heavy development currently, thus its internals are subject to change. However, it could be done with XPflow-specific funcionalities as well. In the long term, using XPflow-specific funcionality would make the framework more generic (less Grid'5000-specific).

4.4. The framework

In this section, we talk about the structure of the implemented framework, its functionalities and usage. The framework consists of two main parts: there is a frontend and there is an XPflow structure called a "library" as a backend. While the expression "library" may sound strange in this context, this is what comes the closest to a class in XPflow: we are able to store variables and have member functions, features which the framework makes much use of.

Below is a more detailed discussion of the most important traits of the frontend and the backend.

4.4.1. Frontend

The frontend is an interface between the user and the framework. It has the function calls the user can use when orchestrating an automated experiment. The example process in

4.2 contains most of the processes that we'll discuss below.

4.4.1.1. init; finish

These processes need to be called at the beginning and at the end of the experiment respectively. These methods have tasks along the lines of variable initialization and metadata collection. They will be discussed in greater detail when talking about the backend.

4.4.1.2. reserve_and_deploy

The allocation and deployment of the nodes are coupled inside one method. The framework currently is only able to allocate nodes on one specified site, but the possibility to run experiments on nodes across multiple sites could be implemented in the future. It's important to note that when using XPflow's *reserve_ nodes*, we set the *keep => true* argument. This is so we actually keep the job open until the given time limit expires - we don't kill it at the end of the experiment. This is so we can use *checkpointing* (see below) to run other experiments on the same nodes without having to do the allocation and deployment process again (which can take up several minutes).

4.4.1.3. broadcast

This method can be used to broadcast any file across the nodes to a specified destination path. It can be utilized to broadcast the benchmark runnable to every node to the destination where the execution will take place. This needs to be done in order for MPI to utilize all of the nodes - if a runnable is not present in the place of execution on a given node, MPI won't be able to start its processes there. For example, if we want to run the experiment and generate the trace files in the /tmp directory, we will broadcast our benchmark runnable to /tmp across the nodes.

4.4.1.4. disable_cores

Disabling cores on nodes can be necessary as a method to overcome simulation inaccuracy caused by SMPI's inability to predict the performance of an MPI application when using multiple cores, as mentioned before, in the *Problem Description* chapter.
This method, unlike the others in the frontend, is an *activity* instead of being a *process*. This is done as a workaround to an XPflow pitfall regarding proxy variables, an issue that will be discussed in 4.4.4.

4.4.1.5. mpirun

This method is for actually running the benchmark itself. It is important to give the full path to the mpirun runnable as an argument to the process call, otherwise we can't be sure that the MPI implementation running is the one we want.

4.4.1.6. trace_gather

After running our MPI experiment, the gathered traces are scattered amongst all the corresponding nodes. This process calls a script called trace_gather[14] (which needs to be set up on the system image manually), which copies the traces collected on the different nodes to the one node we run the script on. We need to call the script from the path where the execution happened, since that's where the traces were gathered. After running this process, we can assume that all the trace files are on the node we made the function call from. For example if we have our traces in /tmp, we need to call trace_gather from there.

Trace_gather has a parameter called arity. This parameter serves the purpose of optimization. Trace_gather works by logging in to each node in question and then sending its respective trace file(s) to the requested destination. If the arity is greater or equal than the number of nodes - 1, then every node sends its package to the head node - if we imagine sending the traces with the help of a tree, we could say that the root is the head node, and all the other nodes are its children. If we specify arity as less than the number of nodes, we can add extra levels to the tree: some nodes don't directly send their traces directly to the head node, but to an intermediate node, which then sends it further up the tree, eventually getting to the root, the head node.

See figure 4.4 for an example: in a system with the total of 8 nodes, we can see on the left how trace_gather works if we set the arity to 7 and when we decrease it to 4. In the first case, all the other 7 nodes send their traces directly to the root of the tree, the head node (H). In the second case, where the arity is 4, we split the nodes into 2 groups: the first group of three have the head node as their parent - they send the traces there directly. The second group can be seen on the right: 3 of the nodes form a third level on the tree, with an intermediate parent. These 3 nodes send their traces to that parent node, after which the parent node sends its traces, coupled with the traces received from its children, to the root node. With proportionally more nodes, not sending everything directly to the root can provide a noticeable performance increase for trace_gather.

Figure 4.4. Examples of the trace_gather's algorithm with two different arity values.

4.4.1.7.   execute_head

This activity can be used to execute any command on the head node. The head node is one of the allocated nodes, arbitrarily chosen by the system. This is the node where the *mpirun* command is run from.

The *execute_head* activity is useful for post-processing purposes, such as calling the TAU[23] script *tau_treemerge.pl*, which is responsible for merging the traces into one trace file and the event files into one event file. This is useful, as after we have one single trace and event file containing all our traces, we can manipulate them easier.

execute_head can also be used to convert our traces. For example, Akypuera's[35] *tau2paje* script can be used to convert TAU traces to traces visualizable with the Pajé[36] visualization tool.

4.4.2.   Backend

As mentioned before, the backend of the framework is an XPflow construct called a *library*. A library has state, so we can have internal variables and methods, a functionality that we make use of: if we couldn't save our state between individual function calls, certain parameters, such as the number of nodes or the name of the site we are working on would need to be given every time. Instead, we can just put those values in variables and use them later when needed.

Libraries can also contain activities. Most of the "real work" is put into the library in the form of member activities. The backend interfaces with the frontend, which makes calls to these activities whenever needed.

4.4.2.1. Metadata

The backend library is also responsible for handling metadata collection. Collected metadata includes:

- the date the experiment was run,

- the runtime of the experiment,

- the benchmark that was run,

- the site and the nodes that were used for the experiment and

- the image that was deployed on the nodes.

Certain elements of this list is provided by the user, while giving certain parameters to a function call. For example when reserving nodes, the user has to give the name of the site he/she wants to use, as well as the name of the image that he/she wants to be deployed. Other metadata is decided during runtime: a prime example to this are the nodes used. The user only tells the framework the number of nodes to allocate - the framework then gives the command on the system to allocate the nodes, then saves the names of the allocated nodes in a nodefile, which can be used later as a reference to our nodes in question. There are also pieces of metadata directly generated by the framework, such as the date and the runtime of the experiment. The date is saved in the previously mentioned *init* method, while the runtime is calculated by saving the starting time in the *init* method, the ending time in the *finish* method, then calculating the difference.

4.4.3. Checkpointing

XPflow has a very useful feature, called checkpointing. A checkpoint can be put anywhere in our experiment, making it possible that in the next run, control is *resumed* at the place of the checkpoint. An example usage can be seen on 4.5: we put a checkpoint after the allocation of nodes (line 12). As previously mentioned in 4.4.1.1, the node allocation is done in such a way that they remain available even after the experiment is done. So with putting a checkpoint there, we don't have to wait through the several minute long node acquisition and deployment process each time we want to run an MPI experiment: we can re-run it on the same allocated nodes, as long as the time limit we set at the node allocation allows us. It is also possible to change the code after the checkpoint to run some other experiment, as it can be seen on 4.6: there, we modify the code so it runs *dt.A.2* instead of *lu.A.2*. We had to modify the code in two places to achieve this: first,

we need to broadcast the other benchmark across the nodes (line 14) and the path to the runnable needs to be changed as well when giving the command to run the experiment (line 19). If the time interval we allocated the nodes for passes, we have to run that part of the experiment again. This can be done by telling XPflow to ignore the checkpoints (with the argument -*I*). In our example, this time interval is 30 minutes, as it can be seen on line 10, where we give it as a parameter to the node acquisition process.

```
1  # -*- coding: utf-8 -*-
2  #!/usr/bin/env xpflow
3  #
4
5  use :g5k
6  require '../test_framework/framework'
7  process :main do
8    run :init, :user => "dlehoczky"
9  #preparation and configuration
10   run :reserve_and_deploy, 2, '120m', 'lille',
11     'wheezy-x64-big-lehoo', '/home/dlehoczky/nodefile_uniq'
12   checkpoint :deployed
13
14   run :broadcast, '/home/dlehoczky/NPB3.3/NPB3.3-MPI/bin/lu.A.2', '/tmp'
15   run :disable_cores
16
17 #the experiment
18   log "mpi experiment starting"
19   run :mpirun, :path => '/tmp/lu.A.2', :n => 2,
20     :outfile => '/tmp/mpi.out'
21
22   run :finish
23 end
24
25 main :main
```

Figure 4.5. Example usage of checkpointing

4.4.4. Proxy variables

There is a shortcoming of XPflow that caused a little headache during the development of the framework. The problem stems from the fact that the XPflow processes are written in a DSL that is embedded in Ruby. Ruby is very flexible in that regard, but one unsolved problem is that the body of the process blocks are executed *before* the execution of the process itself. The problem with this is that at the time of the execution of the process blocks, some of the variables that the process uses contain only some proxy value, not the value that was actually put in there. This problem has been briefly mentioned when talking about the *disable_cores* method. In that particular method, multiple command strings are needed, containing the commands that are needed to be executed to disable the cores on the nodes. In these command strings, the variable containing the nodefile is needed to be inserted. If *disable_cores* was a process, this variable would contain only

```
 1  # -*- coding: utf-8 -*-
 2  #!/usr/bin/env xpflow
 3  #
 4
 5  use :g5k
 6  require '../test_framework/framework'
 7  process :main do
 8    run :init, :user => "dlehoczky"
 9  #preparation and configuration
10    run :reserve_and_deploy, 2, '120m', 'lille',
11      'wheezy-x64-big-lehoo', '/home/dlehoczky/nodefile_uniq'
12    checkpoint :deployed
13
14    run :broadcast, '/home/dlehoczky/NPB3.3/NPB3.3-MPI/bin/dt.A.2', '/tmp'
15    run :disable_cores
16
17  #the experiment
18    log "mpi experiment starting"
19    run :mpirun, :path => '/tmp/dt.A.2', :n => 2,
20      :outfile => '/tmp/mpi.out'
21
22    run :finish
23  end
24
25  main :main
```

Figure 4.6. We can modify the code that comes after the checkpoint

a proxy variable at the time we want to insert it in the string, even if we would put the correct value in it right before. Writing this method as an activity instead of a process solves this problem, since activities are written in plain Ruby.

This chapter was about the implementation details of the framework. First, we talked about the development and testing environment, the Grid'5000 testbed. We went over its architecture and its features deemed to be the most important for this project. After that, we constructed an example process for real-life trace collection, which was used when portraying the features of the experimentation engine, which the framework was developed with: XPflow. We talked about why XPflow is feasible for our project and its main features, through two examples: one basic, general example and one that is very relevant to our case, the aforementioned RL trace collection example. After that part, we talked about the architecture of the implemented framework, and some of its methods. In the end, we mentioned metadata collection as an important feature of the framework, as well as checkpointing, an XPflow feature that proved to be really useful during development.

5. fejezet

Evaluation Plan

In this chapter, we go over the process of the evaluation of the implemented framework, the goal of which is to make conclusions about whether or not we succeeded in achieving the requirements we specified earlier. The evaluation will be done by running a multi-step, fairly typical experiment both manually and with the framework, so we can easily observe the differences.

The experiment process used in the evaluation will consist of the same steps as mentioned before in the Implementation chapter as an example MPI experiment process (see 4.3), used to collect RL traces, which can be used for example for the development of SMPI. The steps in the process are very generic though: we need to be more precise about certain details of the experiment.

5.1. Experiment specification

5.1.1. The benchmark

The benchmark used for evaluation is taken from the NAS Parallel Benchmarks suite[37], which is a small set of benchmarks designed for performance testing of parallel systems. The benchmark chosen for our current purpose is called *lu.B.8*. The name consists of 3 parts. The first part, *lu* indicates what the experiment is about: as its name suggests, the LU benchmark solves a system of equations represented with a matrix, with the LU factorization method.

The second part of the name, *B* is an indication about the complexity of the problem that the benchmark solves. The NPB suite defines so-called "problem classes" for its benchmarks. B is in the middle "standard" complexity category, being more complex

than the Small (S), the Workstation-size (W) and A, the least complex "standard" problem size, but less complex than C and the larger test problem sizes.[38]

Finally, the number *8* at the end of the benchmark's name indicates how many parallel processes it is intended for to be solved: for our purposes, we choose it to be 8, thus, we will allocate 8 nodes for our experiment, each of them running 1 MPI process to solve the problem.

It's important to note that when doing the experiment, we assume that the benchmark is already compiled, in a specified place in the user's home folder: /home/dlehoczky/NPB3.3/NPB3.3-MPI/bin/lu.B.8. For the user *dlehoczky*, the environment variable *$NPB_DIR* is set to /home/dlehoczky/NPB3.3/NPB3.3-MPI, thus it can be used as a shortcut when accessing the benchmark. Also, it is assumed that ssh-keys are set up correctly so the user can reach the specified Grid'5000 site without any password.

5.1.2. The environment

For our experiment, we use the Grid'5000 testbed, discussed in greater detail before, in 4.1. As we said before, there are many different sites to choose from. During development, many tests were needed to be done. Sometimes one or two sites were undergoing maintenance or became unstable for various reasons. Towards the end of the development process, the Lille site proved to be reliable in terms of availability, this is why it was chosen as the site to run our final tests on.

5.1.3. The OS image

We use a customized image called *wheezy-x64-big-lehoo*. This image contains all the necessary tools required to run the described process. These tools were mentioned before, at 3.2.1, when we talked about how to collect RL traces, which this experiment is about. To reiterate, let's sum up what our customized image contains in order to execute our experiment process:

- OpenMPI 1.6.4;

- the *TAU*[23] profiling tool;

- the PAPI[24][25] interface to low-level hardware counters (configured so it's linked to TAU, which uses it for tracing);

- the *Program Database Toolkit (PDT)*[26] (also configured to be linked to TAU);

- the *trace_gather*[14] MPI program;

- *Akypuera*[35], a library to trace mpi applications and generate paje trace files. (In our experiment, we only use its *tau2paje* trace converting script.)

5.1.4. The experiment process

As mentioned before, the experiment process will consist of the same steps as in the example discussed before (see 4.3).

To prepare our experiment process, we start by logging in to the frontend with our user. Then, we start an interactive job, allocating the desired number (in our case, 8) of nodes for a specified amount of time. We won't specify what nodes we want, the system will decide which ones we get. Since there might be performance disparities between two different sets of nodes, we make sure to use the exact same set of nodes in the two experiments (the one done manually and the one done with the framework). We do this by running the experiment with the framework first, then connecting to the job started by the framework and repeating the experiment manually, starting from the deployment part. The node allocation is not really an important part of the experiment, thus, it's not a problem that it's omitted from the manual execution.

As part of the process that we conduct in both experiments, first, we deploy the previously mentioned (5.1.3) customized image on the nodes. Then we broadcast the runnable (*lu.A.8*) to every node's /tmp directory. After that comes the step where we disable all but one core on every node. The reasoning for this step has been discussed previously (see 3.3.1). This concludes the preparation stage.

After that comes the execution of the benchmark. We use OpenMPI 1.6.4, installed on our image.

After the running of the benchmark, comes the post-processing. If we compiled our benchmark correctly with TAU, one trace file (.trc) and one event file (.edf) is generated for each MPI process. The traces are generated on the node of execution. This is why first, we run the *trace_gather*[14] script to collect the traces and event files scattered across all the allocated nodes to the head node. When we have all the files on our head node, we use TAU's *tau_treemerge.pl*, which is a script that merges all our trace files and event files into one trace and event file respectively, also trying to account for any clock skew (see 3.3.3) between the trace files with post-processing methods. Finally, when we have one merged trace and one event file, we can convert our TAU trace to a format that is compatible with the Pajé[36] visualization tool. Visualizing the traces makes it easier to analyze and compare. For the conversion, we use Akypuera's[35] *tau2paje* script.

In this chapter, we went over the specifics of the evaluation of our implementation. First, we talked about the chosen benchmark to gather traces from, which is called *lu.B.8*. We also explained there what the name stands for: LU indicates the problem the

benchmark solves, B is the problem class and 8 indicates how many nodes the benchmark is compiled for. We also discussed what the chosen environment is going to be (Grid'5000) and what operating system image we are going to deploy on our nodes before running the benchmarks on them. Finally, we went over the experiment process: the RL trace collection for a given benchmark. We mentioned that our experiment process will be done both with and without the framework to demonstrate how the two methods differ from each other. We discussed the preparation steps, the running of the benchmark and the post-processing part. Now, in the next chapter, we'll take a look at the results of our experiments.

6. fejezet

Results

In this chapter, we will go over the results of the evaluation, which is done as described in the previous chapter, "Evaluation Plan"5. First comes the part where we discuss the results we got when running the experiment using the framework. Here, we describe how the experiment was orchestrated: what processes were used, how the test file's code looked like. Then we'll take a look at how the experiment was run, what information was given to the user during runtime and how much time it took. Then, we will go over the steps of the experiment done by hand, starting from connecting to the interactive job started by the framework previously, then proceeding with the experiment, describing certain details such as what commands were needed to be given, or how much time the steps took. In the end, we'll compare how the two different experimentation methods fare against each other.

6.1. Experiment using the framework

First, we discuss the experiment orchestrated with the test framework. The test written with the framework consists of one single file, containing 1 method (the main) and 11 function calls. We can look at the code on figure 6.1.
The experiment itself can be divided into three main stages: the preparation/configuration stage, the benchmark stage and the post-processing stage.

```
1  # -*- coding: utf-8 -*-
2  #!/usr/bin/env xpflow
3  #
4
5  use :g5k
6  require '../test_framework/framework'
7  process :main do
8    run :init, :user => "dlehoczky"
9  #preparation and configuration
10   run :reserve_and_deploy, 8, '120m', 'lille',
11     'wheezy-x64-big-lehoo', '/home/dlehoczky/nodefile_uniq'
12   checkpoint :deployed
13
14   run :broadcast, '/home/dlehoczky/NPB3.3/NPB3.3-MPI/bin/lu.B.8', '/tmp'
15
16   run :disable_cores
17 #the experiment
18   log "mpi experiment starting"
19   run :mpirun, :path => '/tmp/lu.B.8', :n => 8,
20     :arguments => '-mca btl ^openib --mca pml ob1 --mca btl_tcp_if_include eth1',
21     :mpirun_path => '/usr/local/openmpi-1.6.4-install/bin/mpirun',
22     :outfile => '/tmp/mpi.out'
23 #post-processing
24   run :trace_gather,
25     :n => 8, :tracegather => '/usr/local/trace_gather/trace_gather',
26     :arity => 4
27   run :run_script,
28     :command => 'tau_treemerge.pl'
29   run :run_script,
30     :command => '/usr/local/akypuera-install/bin/tau2paje',
31     :args => 'tau.trc tau.edf', :outfile => 'lu.B.8.paje', :errorfile => 'tau2paje.error'
32 #metadata
33   run :finish
34 end
35 main :main
```

Figure 6.1. The experiment, orchestrated with the test framework.

6.1.1. Preparation

6.1.1.1. Experiment initialization

In the beginning of every experiment, the *init* method has to be called in order to configure certain variables with default values and also to make configuration steps regarding metadata collection (e.g. record the date of the experiment, the exact time of starting, etc.). Also, here is where the user can set his/her username to log in to the nodes. If not set, the username will default to root.

6.1.1.2. Reserve and deploy

This is the part where the program starts a job by allocating 8 nodes on the *lille* site and deploying the image (given as a parameter to the method) on them. We set the

allocation time to 120 minutes (2 hours), which should give us plenty of time to execute the experiment with both the framework and manually on the nodes.

As part of the process, a so-called 'nodefile' is created, which contains the names of all the nodes that we allocated. Originally, the environment variable *$OAR_ FILE_ NODES* is set on the frontend when logged in to an active job, with the value of the path to a file containing the node names, each as many times as many cores they have. During the experiment, we create a nodefile that only contains each node once, so it can be used as a machinefile for certain commands later. We specify a custom path to the nodefile with an argument. If that argument is not set, our nodefile would be created in the home folder anyway.

6.1.1.3. Broadcasting the runnable

The benchmark needs to be copied to every node to the path where the execution will take place - otherwise MPI wouldn't be able to start its processes on every node. The parameters here are pretty straightforward - we give the method the path to the compiled benchmark and the path to copy it to. The nodefile is not needed to be given, since it's already been created and stored by the library before.

6.1.1.4. Disabling cores

This is the method call that is responsible for disabling all but one core on every node that takes part in the execution. We talked about why this is necessary earlier, see 3.3.1. No parameters are needed to be given here.

6.1.2. Running the benchmark

After all the configuration steps have been successfully done, it's time for actually running the benchmark. As parameters, we first give the *mpirun* method call the path to the runnable (meaning, of course, the path that we broadcasted the runnable to before, not its original path) and the number of nodes, which, in our case, is 8. We also specify the MPI runnable with its full path: this is to avoid confusion about which version of MPI is actually running on the nodes, ensuring that it's the version we deployed on the image (OpenMPI 1.6.4). If not given, it would be defaulted to simply 'mpirun'. We specify the file to put the output of the benchmark to as well.

We can put any extra arguments that we want to give to mpirun in the "arguments" parameter. In this case, we use this parameter to specify that our benchmark must not use Infiniband. This is needed because SMPI is unable to simulate Infiniband connection,

thus, a trace acquired using it would show significantly better performance than the simulated version.

6.1.3. Post-processing

6.1.3.1. Gathering the traces

As we mentioned before, in 5.1.4, since we compiled our benchmark with TAU, every MPI process produces one trace (.trc) and one event (.evt) file. These traces are scattered across the nodes. To gather them to one node, we use the *trace_ gather* MPI program. The method call with the same name takes as parameter the number of nodes, the path to the trace_gather executable (the MPI executable is remembered from the *mpirun* call) and the arity parameter (see 4.4.1.6).

6.1.3.2. Merging the traces

For running *tau_ treemerge.pl*, we use the *run_ script* method of the framework, which is a generic method that executes the command given as parameter on the head node. No other arguments are needed in this case. Although we could specify files to put the output and the error messages to, we don't really need that in this simple case. The directory to merge the traces in defaults to "/tmp", so that is not needed either. As we mentioned before, this script doesn't only merge our traces, it also attempts at accounting for any clock skew.

6.1.3.3. Converting the traces

As mentioned before, we use Akypuera's[35] *tau2paje* script to convert our traces to Pajé-visualizable format. For this, we use the *run_ script* method again, giving it the trace and event files as arguments, as well as files to put the output and the error messages to. Tau2paje generates error messages because of synchronization problems (clock skew). As we described before (see 3.3.3), clock skew problems are really hard to overcome. Although the merging script before tried to take care of the problem, it can't eliminate it completely, but it does a well enough job so this problem doesn't interfere with the conversion results in any noticeable way.

After the conversion is done, the visualizable trace file called *lu.B.8.paje* is available on the /tmp folder of the head node. We can now copy it to our local machine and visualize the result. (The results will be discussed below, after discussing the manual experiment.)

6.1.4. Metadata

In the end, we get some amount of metadata generated by the experiment, summarizing some of its most important specifics, as it can be seen on figure ref.
#screenshot of the metadata

6.2. Experiment by hand

Now, let's take a look at how the manual experiment went. As mentioned before, the two experiments consist of the same exact steps. As before, we divide the experiment into three main sections: preparation stage, benchmark stage and post-processing stage.

6.2.1. Preparation

Below is a table summarizing the results for the preparation and configuration stage. The whole preparation stage took place on the frontend of the site "lille", with the user "dlehoczky".

Preparation		
Task	*Command*	*Time (s)*
Reserving nodes	Omitted - joined the existing job, using the already allocated nodes	
Deployment	`kadeploy3 -f $OAR_FILE_NODES -e wheezy-x64-big-lehoo`	
Create a nodefile containing all the node names only once	`cat $OAR_FILE_NODES	uniq > ~/nodefile_uniq`	
Broadcast of runnable	`taktuk -l root -f ~/nodefile_uniq broadcast put  $NPB_DIR/bin/lu.B.8   /tmp`	
Disabling cores	`taktuk -l root -f /home/dlehoczky/nodefile_uniq broadcast exec [ 'for i in /sys/devices/system/cpu/ cpu[1-9]*/online; do echo 0 > "$i" ; done' ]`	
Optional: Check if the cores have been disabled accordingly	`taktuk -f ~/nodefile_uniq broadcast exec ['cat /proc/cpuinfo']	grep process	awk 'if($9>0) print $1'	uniq	awk -F".fr-" 'print $1".fr"'`	

As a side note: when doing the core disabling with the framework, the program repeats the disabling step until it confirms it to be successful, with the step marked as "Optional" here. In the step where we check if the cores are disabled, every node name is displayed that has more than 1 core running. The disabling step needs to be repeated until we don't see any node names when running that command. When running the test, the disabling step was successful the first time, thus there was no need to repeat that step.

6.2.2. Running the benchmark

Below is the table about the command that was given to run the benchmark (which is exactly the same as it was for the framework), as well as its runtime. The benchmark had to be run from the directory the benchmark was broadcasted to (/tmp in our case), on the designated head node (which was, in our case, #TODO: head node, the same as it was for the framework.

Running the benchmark		
Task	*Command*	*Time (s)*

Running the benchmark	`/usr/local/openmpi-1.6.4-install`
	`/bin/mpirun -mca btl ^openib -mca pml`
	`ob1 -machinefile ~/nodefile_uniq -np 8`
	`/tmp/lu.B.8 1>/tmp/mpi.out`	

6.2.3. Post-processing

After executing the benchmark, the produced traces are sitting on their respective nodes. Below, we summarize how the post-processing of the files went. This stage of the experiment, as it was for the benchmark, was executed on the head node, from the /tmp directory.

Post-processing		
Task	*Command*	*Time (s)*

Gathering the traces	`/usr/local/bin/mpirun -machinefile`
	`~/nodefile_uniq -np 8`
	`/usr/local/trace_gather/trace_gather -f`
	`1 -a 4 -m ~/nodefile_uniq`	

Merging the traces	`tau_treemerge.pl`	

Converting the trace to Pajé-compatible format	`/usr/local/akypuera-install/bin/tau2paje`
	`tau.trc tau.edf 1>lu.B.8.paje`
	`2>tau2paje.error`	

As mentioned previously at the end of the experiment with the framework: after the conversion process, the file *lu.B.8.paje* is available on the head node, in its /tmp folder. Now, we can use scp, rcp or some other program to download it to our local machine from the remote node.

6.3. Comparison

First, let's take a look at how much time the experiments took. Below is a table summarizing the elapsed time for both experiments.

Stage	Time (s)	
	Experiment with the framework	*Manual experiment*
Preparation	-	-
Running the benchmark	-	-
Post-processing	-	-

Now, let's take a look at the generated traces. Below are the results for both experiments, visualized with Vite, a Pajé visualization tool. #TODO: visualizations here; below: compare them advantages of the framework (obvious one: no user intervention required during the process)

7. fejezet

Conclusion

A. Függelék

Appendix Title Here

Write your Appendix content here.

Bibliography

[1] Message Passing Interface, Forum. *MPI: A Message-Passing Interface Standard Version 3.0*, September 2012.

[2] Gabriel, E., Fagg G., E., Bosilca, G., Angskun, T., Dongarra J., J., Squyres J., M., Sahay, V., Kambadur, P., Barrett, B., Lumsdaine, A., Castain R., H., Daniel D., J., Graham R., L., and Woodall T., S. Open mpi: Goals and concept, and design of a next generation mpi implementation. In *Proceedings, 11th European PVM/MPI Users' Group Meeting*, pages 97–104, Budapest, Hungary, September 2004.

[3] , MPICH. Mpich overview, 2012. URL http://www.mpich.org/about/overview/. Last visited on 07/08/2013.

[4] Casanova, H., Legrand, A., and Quinson, M. Simgrid: a generic framework for large-scale distributed experiments. In *Proceedings of the Tenth International Event on Computer Modeling and Simulation*, pages 126–131, 2008.

[5] Clauss, P.-N., Stillwell, M., Genaud, S., Suter, F., Casanova, H., and Quinson, M. Single node on-line simulation of mpi applications with smpi. In *International Parallel & Distributed Processing Symposium*, pages 664–675, May 2011.

[6] Mesirov J., P. Accessible reproducible research. *Science Magazine*, 327: 415–416, 2010. URL http://statlab.bio5.org/foswiki/pub/Main/PapersForClassCPH685/science-reproducible-research.pdf.

[7] Buchert, T. and Nussbaum, L. Using business workflows to improve quality of experiments in distributed systems research. In *SC12 - SuperComputing 2012 (poster session)*, Salt Lake City, United States, August 2012. URL http://hal.inria.fr/hal-00724312.

[8] Thakur, R., Rabenseifner, R., and Gropp, W. Optimization of collective communication operations in mpich. *Int'l Journal of High Performance Computing Applications*, pages 49–66, 2005.

[9] Buntinas D. Mercier, D. and Gropp, W. Implementation and evaluation of shared-memory communication and synchronization operations in mpich2 using the nemesis communication subsystem. *Parallel Computing*, 33:634–644, September 2007. URL http://linkinghub.elsevier.com/retrieve/pii/S0167819107000786.

[10] Hillston, J. Performance modelling and lecture 1 ("modelling and simulation"), 2012. URL http://www.inf.ed.ac.uk/teaching/courses/pm/PM-lecture1.pdf.

[11] Hillston, J. Performance modelling and lecture 13 ("simulation models: Introduction and motivation"), 2012. URL http://www.inf.ed.ac.uk/teaching/courses/pm/PM-lecture1.pdf.

[12] Bédaride, P., Degomme, A., Genaud, S., Legrand, A., Markomanolis G., S., Quinson, M., Stillwell, M., Suter, F., and Videau, B. Improving simulations of mpi applications using a hybrid network model with topology and contention support. 2013.

[13] Desprez, F., Markomanolis G., S., Suter, F., and Quinson, M. Assessing the performance of mpi applications through time-independent trace replay. In *Proceedings of the 2nd International Workshop of Parallel Software Tools and Tool Infrastructures (PSTI)*, pages 467–476, Taipei, Taiwan, September 2011.

[14] Markomanolis G., S. and Suter, F. Time-independent trace acquisition framework – a grid'5000 how-to. Technical report, Institut National de Recherche en Informatique et en Automatique, 2011. URL http://hal.inria.fr/inria-00593842.

[15] Videau, B. and Richard, O. Expo : un moteur de conduite d'expériences pour plates-forme dédiées. In *Conférence Française en Systèmes d'Exploitation (CFSE)*, 2008.

[16] Rakotoarivelo, T., Ott, M., Seskar, I., and Jourjon, G. Omf: a control and management framework for networking testbeds. In *SOSP Workshop on Real Overlays and Distributed Systems (ROADS '09)*, page 6, Big Sky, USA, October 2009.

[17] Deelman, E., Singh, G., Si, M.-H., Blythe, J., Gil, Y., Kesselman, C., Mehta, G., Vahi, K., Berriman B., G., Good, J., Laity, A., Jacob J., C., and Katz D., S. Pegasus: a framework for mapping complex scientific workflows onto distributed systems. *Scientific Programming Journal*, 13:219–237, 2005.

[18] Altintas, I., Berkley, C., Jaeger, E., Jones, M., Ludascher, B., and Mock, S. Kepler: An extensible system for design and execution of scientific workflows. In *16th International Conference on Scientific and Statistical Database Management*, Santorini Island, Greece, June 2004.

[19] Wolstencroft, K., Haines, R., Fellows, D., Williams, A., Withers, D., Owen, S., Soiland-Reyes, S., Dunlop, I., Nenadic, A., Fisher, P., Bhagat, J., Belhajjame, K., Bacall, F., Hardisty, A., Nieva de la, Hidalga A., Vargas M. P., B., Sufi, S., and Goble, C. The taverna workflow suite: designing and executing workflows of web services on the desktop, web or in the cloud. *Nucleic Acids Research*, 41:557–561, May 2013.

[20] Sonntag, M., Karastoyanova, D., and Deelman, E. Bridging the gap between business and scientific workflows: Humans in the loop of scientific workflows. In *e-Science (e-Science), 2010 IEEE Sixth International Conference on*, pages 206–213. IEEE, 2010.

[21] Bechhofer, S., Ainsworth, J., Bhagat, J., Buchan, I., Couch, P., Cruickshank, D., Roure D., D., Delderfield, M., Dunlop, I., and Gamble, M. Why linked data is not enough for scientists. In *e-Science and 2010 IEEE Sixth International Conference*, pages 300–307, 2010. URL http://ieeexplore.ieee.org/xpls/abs_all.jsp?arnumber=5693931.

[22] Buchert, T. and Nussbaum, L. Leveraging business workflows in distributed systems research for the orchestration of reproducible and scalable experiments. In *MajecSTIC 2012*, Lille, France, October 2012.

[23] Shende S., S. and Malony A., D. The tau parallel performance system. *International Journal of High Performance Computing Applications*, 20:287–311, 2006. URL http://statlab.bio5.org/foswiki/pub/Main/PapersForClassCPH685/science-reproducible-research.pdf.

[24] Mucci P., J., Browne, S., Deane, C., and Ho, G. Papi: A portable interface to hardware performance counters. In *Proceedings of Department of Defense HPCMP Users Group Conference*, California, USA, June 1999.

[25] London, K., Moore, S., Mucci P., J., Seymour, K., and Luczak, R. The papi cross-plaftorm interface to hardware performance counters. In *Department of Defense Users' Group Conference Proceedings*, Biloxi, Mississippi, USA, June 2001.

[26] Lindlan K., A., Cuny, J., Malony A., D., Shende, S., Mohr, B., Rivenburgh, R., and Rasmussen, C. A tool framework for static and dynamic analysis of object-oriented software with templates. In *Supercomputing, ACM/IEEE 2000 Conference*, Dallas, Texas, USA, 2000.

[27] Desprez, F., Markomanolis G., S., and Suter, F. Improving the accuracy and efficiency of time-independent trace replay. Technical report, Institut National de Recherche en Informatique et en Automatique, October 2012.

[28] Wu. C., E., Bolmarcich, A., Snir, M., Wootton, D., Parpia, F., Chan, A., Lusk, E., and Gropp, W. From trace generation to visualization: A performance framework for distributed parallel systems. In *Supercomputing, ACM/IEEE 2000 Conference*, pages 50–67, Dallas, Texas, USA, 2000.

[29] Becker, D., Rabenseifner, R., Wolf, F., and Linford J., C. Scalable timestamp synchronization for event traces of message-passing applications. *Parallel Computing*, 35: 595–607, 2009.

[30] Mills D., L. Network time protocol (version 3). Technical report, The Internet Engineering Task Force - Network Working Group, March 1992.

[31] Bolze, R., Cappello, F., Caron, E., Daydé, M., Desprez, F., Jeannot, E., Jégou, Y., Lanteri, S., Leduc, J., Melab, N., Mornet, G., Namyst, R., Primet, P., Quetier, B., Richard, O., Talbi, E., and Touche, I. Grid'5000: A large scale and highly reconfigurable experimental grid testbed. *International Journal of High Performance Computing Applications*, 20:481–494, November 2006.

[32] Capit, N., Da Costa, G., Georgiou, Y., Huard, G., Martin, C., Mounié, G., Neyron, P., and Richard, O. A batch scheduler with high level components. In *IEEE International Symposium on Cluster Computing and the Grid*, volume 2, Cardiff, Wales, UK, May 2005.

[33] Claudel, B., Huard, G., and Richard, O. Taktuk, adaptive deployment of remote executions. In *Proceedings of the International Symposium of High Performance Distributed Computing (HPDC)*, Munich, Germany, June 2009.

[34] Jeanvoine, E., Sarzyniec, L., and Nussbaum, L. Kadeploy3: Efficient and scalable operating system provisioning. *Usenix*, 38:38–44, February 2013.

[35] Schnorr, L. Akypuera, a library to trace mpi applications and generate paje trace files, 2013. URL https://github.com/schnorr/akypuera. Last visited on 08/08/2013.

[36] Chassin de, Kergommeaux D., de, Oliveira Stein B., and Bernard P., E. Performance characteristics of the multi-zone nas parallel benchmarks. *Parallel Computing*, 26: 1253–1274, August 2000.

[37] Jin, H., Frumkin, M., and Yan, J. The openmp implementation of nas parallel benchmarks and its performance. Technical report, Technical Report NAS-99-011, NASA Ames Research Center, October 1999.

[38] Dunbar, J. Problem sizes and parameters in nas parallel benchmarks, 2013. URL http://www.nas.nasa.gov/publications/npb_problem_sizes.html. Last visited on 07/08/2013.